

CS101 Algorithms and Data Structures

Trees

Textbook Ch B.5, 10.4


Outline

- Tree structure
- Implementation
- Tree traversal
- Forest

Trees

A rooted tree data structure stores information in *nodes*

- There is a first node, or *root*
- Each node has variable number of references to successors
- Each node, other than the root, has exactly one node pointing to it




Terminology

All nodes will have zero or more child nodes or *children*

- I has three children: J, K and L

For all nodes other than the root node, there is one parent node

- H is the parent of I




Terminology

The *degree* of a node is defined as the number of its children:

$$\text{deg}(I) = 3$$

Nodes with the same parent are *siblings*


- J, K, and L are siblings



Terminology


Nodes with degree zero are also called *leaf nodes*

All other nodes are said to be *internal nodes*, that is, they are internal to the tree



Terminology


Leaf nodes:



Wesley-Hunt, G. D.; Flynn, J. J. "Phylogeny of the Carnivora: basal relationships among the Carnivoramorphans, and assessment of the position of 'Miacoidea'"

Terminology


Internal nodes:



Wesley-Hunt, G. D.; Flynn, J. J. "Phylogeny of the Carnivora: basal relationships among the Carnivoramorphans, and assessment of the position of 'Miacoidea'"

Terminology

These trees are equal if the order of the children is ignored (*unordered trees*)



They are different if order is relevant (*ordered trees*)

- We will usually examine ordered trees (linear orders)


Terminology

A **path** is a sequence of nodes: (a_0, a_1, \dots, a_n)

where a_{k+1} is a child of a_k


The length of this path is n

e.g., the path (B, E, G)
has length 2



Terminology

Paths of length 10 (11 nodes) and 4 (5 nodes)




Wesley-Hunt, G. D.; Flynn, J. J. "Phylogeny of the Carnivora: basal relationships among the Carnivoramorpha, and assessment of the position of 'Miacoidea'"

Terminology

For each node in a tree, there exists a unique path from the root node to that node


The length of this path is the *depth* of the node, e.g.,

- E has depth 2
- L has depth 3



Terminology

Nodes of depth up to 17



Wesley-Hunt, G. D.; Flynn, J. J. "Phylogeny of the Carnivora: basal relationships among the Carnivoramorphan, and assessment of the position of 'Miacoidea'"

Terminology

The *height* of a tree is defined as the maximum depth of any node within the tree


The height of a tree with one node is 0

- Just the root node

For convenience, we define the height of the empty tree to be -1

Terminology

The height of this tree is 17



Wesley-Hunt, G. D.; Flynn, J. J. "Phylogeny of the Carnivora: basal relationships among the Carnivoramorphan, and assessment of the position of 'Miacoidea'"

Terminology

If a path exists from node a to node b :

- a is an *ancestor* of b
- b is a *descendant* of a


Thus, a node is both an ancestor and a descendant of itself

- We can add the adjective *strict* to exclude equality: a is a *strict descendant* of b if a is a descendant of b but $a \neq b$


The root node is an ancestor of all nodes

Terminology

The descendants of node B are B, C, D, E, F, and G:




The ancestors of node I are I, H, and A:



Terminology


All descendants (including itself) of the indicated node



Wesley-Hunt, G. D.; Flynn, J. J. "Phylogeny of the Carnivora: basal relationships among the Carnivoramorphans, and assessment of the position of 'Miacoidea'"

Terminology

All ancestors (including itself) of the indicated node




Wesley-Hunt, G. D.; Flynn, J. J. "Phylogeny of the Carnivora: basal relationships among the Carnivoramorphan, and assessment of the position of 'Miacoidea'"

Terminology

A **recursive definition** of a tree:

- A degree-0 node is a tree
- A node with degree n is a tree if it has n children and all of its children are disjoint trees (i.e., with no intersecting nodes)

Given any node a within a tree, the collection of a and all of its descendants is said to be a *subtree of the tree with root a*



Outline

- Tree structure
- **Implementation**
- Tree traversal
- Forest

Operations

Operations on a tree include:

- Accessing the root:
- Given an object in the container:
 - Access the parent of the current object
 - Find the degree of the current object
 - Get a reference to a child,
 - Attach a new sub-tree to the current object
 - Detach this tree from its parent


Implementation

We can implement a general tree by using a class which:

- Stores an element
- Stores the children in a list

Implementation

The tree with six nodes would be stored as follows:



Implementation

Suppose we want to find the **size** of a tree:

- If there are no children, the size is 1
- Otherwise, the size is one plus the size of all the children

```
template <typename Type>
int Simple_tree<Type>::size() const {
    int s = 1;

    for (
        Single_node<Simple_tree *> *ptr = children.head();
        ptr != nullptr;
        ptr = ptr->next()
    ) {
        s += ptr->retrieve()->size();
    }

    return s;
}
```

Implementation

Suppose we want to find the **height** of a tree:

- If there are no children, the height is 0
- Otherwise, the height is one plus the maximum height of any sub tree


```
template <typename Type>
int Simple_tree<Type>::height() const {
    int h = 0;

    for (
        Single_node<Simple_tree *> *ptr = children.head();
        ptr != nullptr;
        ptr = ptr->next()
    ) {
        h = std::max( h, 1 + ptr->retrieve()->height() );
    }


    return h;
}
```

Alternative Implementation

Each node has a pointer to a list containing its children.




Each node has 2 pointers: one to its first child and one to next sibling




Parent Pointer Tree

A parent pointer tree is a tree where each node only keeps a reference to its parent node





Parent Pointer Tree

This requires significantly less memory than our general tree structure, as no data structure is required to track the children



Implementation

Instead, generate an array of size n and associate each entry with a node in the tree




Implementation

Store the index of the parent in each node

- The root node, wherever it is, points to itself

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19
0	0	0	0	0	2	2	2	3	3	4	4	4	4	5	5	10	12	12	15
A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T



Outline

- Tree structure
- Implementation
- Tree traversal
- Forest

Tree Traversals

Question: how can we iterate through all the objects in a tree in a predictable and efficient manner

- Requirements: $\Theta(n)$ run time and $o(n)$ memory


Two types of traversals

- Breadth-first traversal
- Depth-first traversal

Breadth-First Traversal

Breadth-first traversals visit all nodes at a given depth before descending a level

- Order: A B H C D G I E F J K



Breadth-First Traversal


The easiest implementation is to use a queue:

- Place the root node into a queue
- While the queue is not empty:
 - Pop the node at the front of the queue
 - Push all of its children into the queue

The order in which the nodes come out of the queue will be in breadth-first order


Breadth-First Traversal

Push the root directory A




Breadth-First Traversal

Pop A and push its two sub-directories: B and H




Breadth-First Traversal

Pop B and push C, D, and G




Breadth-First Traversal

Pop H and push its one sub-directory I




Breadth-First Traversal

Pop J



Breadth-First Traversal

Pop K and the queue is empty




Breadth-First Traversal

The resulting order

A B H C D G I E F J K

is in breadth-first order:



Breadth-First Traversal

Computational complexity


- Run time is $\Theta(n)$
- Space: maximum nodes at a given depth, $O(n)$

Depth-first Traversal

A backtracking algorithm for stepping through a tree:

- At any node, proceed to the first child that has not yet been visited
- If we have visited all the children (of which a leaf node is a special case), backtrack to the parent and repeat this process


We end once all the children of the root are visited



Depth-first Traversal

Each node is visited multiple times in such a scheme


- First time: before any children
- Last time: after all children, before backtracking



Pre-ordering

Ordering nodes by their first visits results in the sequence:

A, B, C, D, E, F, G, H, I, J, K, L, M



Implementing Depth-First Traversals

Implementation with recursion:

```
template <typename Type>
void Simple_tree<Type>::depth_first_traversal() const {
    // Perform pre-visit operations on the element
    std::cout << element << ' ';

    // Perform a depth-first traversal on each of the children
    for (
        Single_node<Simple_tree *> *ptr = children.head();
        ptr != 0; ptr = ptr->next()
    ) {
        ptr->retrieve()->depth_first_traversal();
    }


    // Perform post-visit operations on the element
    std::cout << element << ' ';
}
```

Implementing Depth-First Traversals


Alternatively, we can use a stack:

- Create a stack and push the root node onto the stack
- While the stack is not empty:
 - Pop the top node
 - Push all of the children of that node to the top of the stack **in reverse order**

DFS using a Stack




DFS using a Stack



D
E
H

DFS using a Stack



E
H

Implementing Depth-First Traversals

Computational complexity of DFS using stack

- Run time is $\Theta(n)$
- The objects on the stack are all unvisited siblings from the root to the current node
 - If each node has a maximum of two children, the memory required is $\Theta(h)$: the height of the tree

Computational complexity of DFS using recursion?

- The same complexity?

Guidelines

DFS is used when at each node we need information from:

- all its children or descendants, or
- all its ancestors

In designing a DFS, it is necessary to consider:

1. Before the children are traversed, what initializations, operations and calculations must be performed?
2. In recursively traversing the children:
 - a) What information must be passed to the children during the recursive call?
 - b) What information must the children pass back, and how must this information be collated?
3. Once all children have been traversed, what operations and calculations depend on information collated during the recursive traversals?
4. What information must be passed back to the parent?


Applications of DFS

Displaying information about directory structures and the files contained within

- Printing a hierarchical structure
- Determining memory usage

Printing a Hierarchy

Consider the directory structure presented on the left—how do we display this in the format on the right?



What do we do at each step?

Printing a Hierarchy


Assume the function `void print_tabs(int n)` prints n tabs

```
template <typename Type>
void Simple_tree<Type>::print( int depth ) const {
    print_tabs( depth );
    std::cout << retrieve()->name() << '/' << std::endl;

    for (
        Single_node<Simple_tree*> *ptr = children.head();
        ptr != 0; ptr = ptr->next()
    ) {
        ptr->retrieve()->print( depth + 1 );
    }
}
```

Determining Memory Usage

Printing the directory memory usage of this directory structure:



```
bin/ 12
local/ 15
usr/ 31
adm/ 6
cron/ 5
log/ 9
var/ 23
/ 61
```

Printing a Hierarchy

```
template <typename Type>
int Simple_tree<Type>::du( int depth ) const {
    int usage = retrieve()->memory();

    for (
        Single_node<Simple_tree *> *ptr = children.head();
        ptr != 0; ptr = ptr->next()
    ) {
        usage += ptr->retrieve()->du( depth + 1 );
    }

    print_tabs( depth );
    std::cout << retrieve()->name() << "/" << usage << std::endl;


    return usage;
}
```

Height

The `int height()` const function is recursive in nature:

1. Before the children are traversed, we assume that the node has no children and we set the height to zero: $h_{\text{current}} = 0$
2. In recursively traversing the children, each child returns its height h and we update the height if $1 + h > h_{\text{current}}$
3. Once all children have been traversed, we return h_{current}

When the root returns a value, that is the height of the tree




Outline

- Tree structure
- Implementation
- Tree traversal
- Forest

Forest

A rooted **forest** is a data structure that is a collection of disjoint rooted trees




Forest

Note that:

- Any tree can be converted into a forest by removing the root node
- Any forest can be converted into a tree by adding a root node that has the roots of all the trees in the forest as children


Traversals

Traversals on forests can be achieved by treating the roots as children of a notional root



Traversals

Traversals on forests can be achieved by treating the roots as children of a notional root




Pre-order traversal: A E B O F P U G H C I Q V W J D K R L M S T X Y Z N

Post-order traversal: E A O U P F G H B V W Q I J C R K L S X Y Z T M N D

Breadth-first traversal: A B C D E F G H I J K L M N O P Q R S T U V W X Y Z

Application

- In Windows, each drive forms the root of its own directory structure
- Each of the directories is hierarchical—that is, a rooted tree



Summary

- Tree structure
 - Terminology
- Implementation
 - Children in a list
- Tree traversal
 - BFS, using queue
 - DFS, using recursion or stack
- Forest